

Data preprocessing

1.Import the Libraries

```
In [1]:
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

2.Importing the dataset.

```
In [2]:
dataset = pd.read_csv("WA_Fn-UseC_-HR-Employee-Attrition.csv")
dataset.head(5)
```

Out[2]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	...	Rela
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences1	1	1	...	
1	49	No	Travel_Frequently	279	8	Research & Development	1	Life Sciences1	2	...		
2	37	Yes	Travel_Rarely	1373	2	Research & Development	2	Other	1	4	...	
3	33	No	Travel_Frequently	1392	3	Research & Development	4	Life Sciences1	5	...		
4	27	No	Travel_Rarely	591	2	Research & Development	1	Medical	1	7	...	

rows × 35 columns

```
In [3]:
dataset.shape
```

Out[3]:(1470, 35)

3.Checking for Null Values.

```
In [4]:
dataset.isnull().any()
```

Out[4]:

Age	False
Attrition	False
BusinessTravel	False
DailyRate	False
Department	False
DistanceFromHome	False
Education	False
EducationField	False
EmployeeCount	False
EmployeeNumber	False
EnvironmentSatisfaction	False
Gender	False
HourlyRate	False
JobInvolvement	False
JobLevel	False
JobRole	False
JobSatisfaction	False
MaritalStatus	False
MonthlyIncome	False
MonthlyRate	False
NumCompaniesWorked	False
Over18	False
OverTime	False
PercentSalaryHike	False
PerformanceRating	False

```
RelationshipSatisfaction  False
StandardHours            False
StockOptionLevel         False
TotalWorkingYears        False
TrainingTimesLastYear    False
WorkLifeBalance          False
YearsAtCompany           False
YearsInCurrentRole       False
YearsSinceLastPromotion  False
YearsWithCurrManager     False dtype:
bool
```

```
In [5]:
dataset.isnull().sum()
```

```
Out[5]:Age          0
Attrition          0
BusinessTravel     0
DailyRate         0
Department        0
DistanceFromHome   0
Education         0
EducationField     0
EmployeeCount      0
EmployeeNumber     0
EnvironmentSatisfaction  0
Gender            0
HourlyRate        0
JobInvolvement    0
JobLevel         0
JobRole          0
JobSatisfaction   0
MaritalStatus     0
MonthlyIncome     0
MonthlyRate       0
NumCompaniesWorked  0
Over18           0
OverTime          0
PercentSalaryHike  0
PerformanceRating  0
RelationshipSatisfaction  0
StandardHours     0
StockOptionLevel  0
TotalWorkingYears  0
TrainingTimesLastYear  0
WorkLifeBalance   0
YearsAtCompany    0
YearsInCurrentRole  0
YearsSinceLastPromotion  0
YearsWithCurrManager  0 dtype:
int64
```

```
In [6]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Age                 1470 non-null  int64
1   Attrition           1470 non-null  int64
2   BusinessTravel      1470 non-null  object
3   DailyRate           1470 non-null  int64
4   Department          1470 non-null  object
5   DistanceFromHome    1470 non-null  int64
6   Education            1470 non-null  int64
7   EducationField      1470 non-null  object
8   EmployeeCount       1470 non-null  int64
9   EmployeeNumber      1470 non-null  int64
10  EnvironmentSatisfaction  1470 non-null  int64
11  Gender              1470 non-null  object
12  HourlyRate          1470 non-null  int64
13  JobInvolvement      1470 non-null  int64
14  JobLevel            1470 non-null  int64
15  JobRole             1470 non-null  object
16  JobSatisfaction     1470 non-null  int64
17  MaritalStatus       1470 non-null  object
18  MonthlyIncome       1470 non-null  int64
```

```
19 MonthlyRate 1470 non-null int64
20 NumCompaniesWorked 1470 non-null int64
21 Over18 1470 non-null object
22 OverTime 1470 non-null object
23 PercentSalaryHike 1470 non-null int64
24 PerformanceRating 1470 non-null int64
25 RelationshipSatisfaction 1470 non-null int64
26 StandardHours 1470 non-null int64
27 StockOptionLevel 1470 non-null int64
28 TotalWorkingYears 1470 non-null int64
29 TrainingTimesLastYear 1470 non-null int64
30 WorkLifeBalance 1470 non-null int64
31 YearsAtCompany 1470 non-null int64
32 YearsInCurrentRole 1470 non-null int64
33 YearsSinceLastPromotion 1470 non-null int64 34 YearsWithCurrManager 1470 non-null int64 dtypes: int64(27), object(8) memory usage: 402.1+ KB
In [4]:
```

dataset['Attrition'] = dataset['Attrition'].map({'Yes': 1, 'No': 0})

In [13]: dataset.head(5)

Out[13]:

	Age	Attrition	BusinessTravel	DailyRate	Research &	Sales	1	2	Life Sciences	1	1	...
0	41	1	Travel_Rarely	1102	Development	8		1	Life Sciences	1		2 ...
1	49	0	Travel_Frequently	279	Research &							
2	37	1	Travel_Rarely	1373	Development				2	2 Other	1	4 ...
3	33	0	Travel_Frequently	1392	Research &			3	4 Life Sciences	1	5	...
4	27	0	Travel_Rarely	591	Development	2	1		Medical	1		7 ...

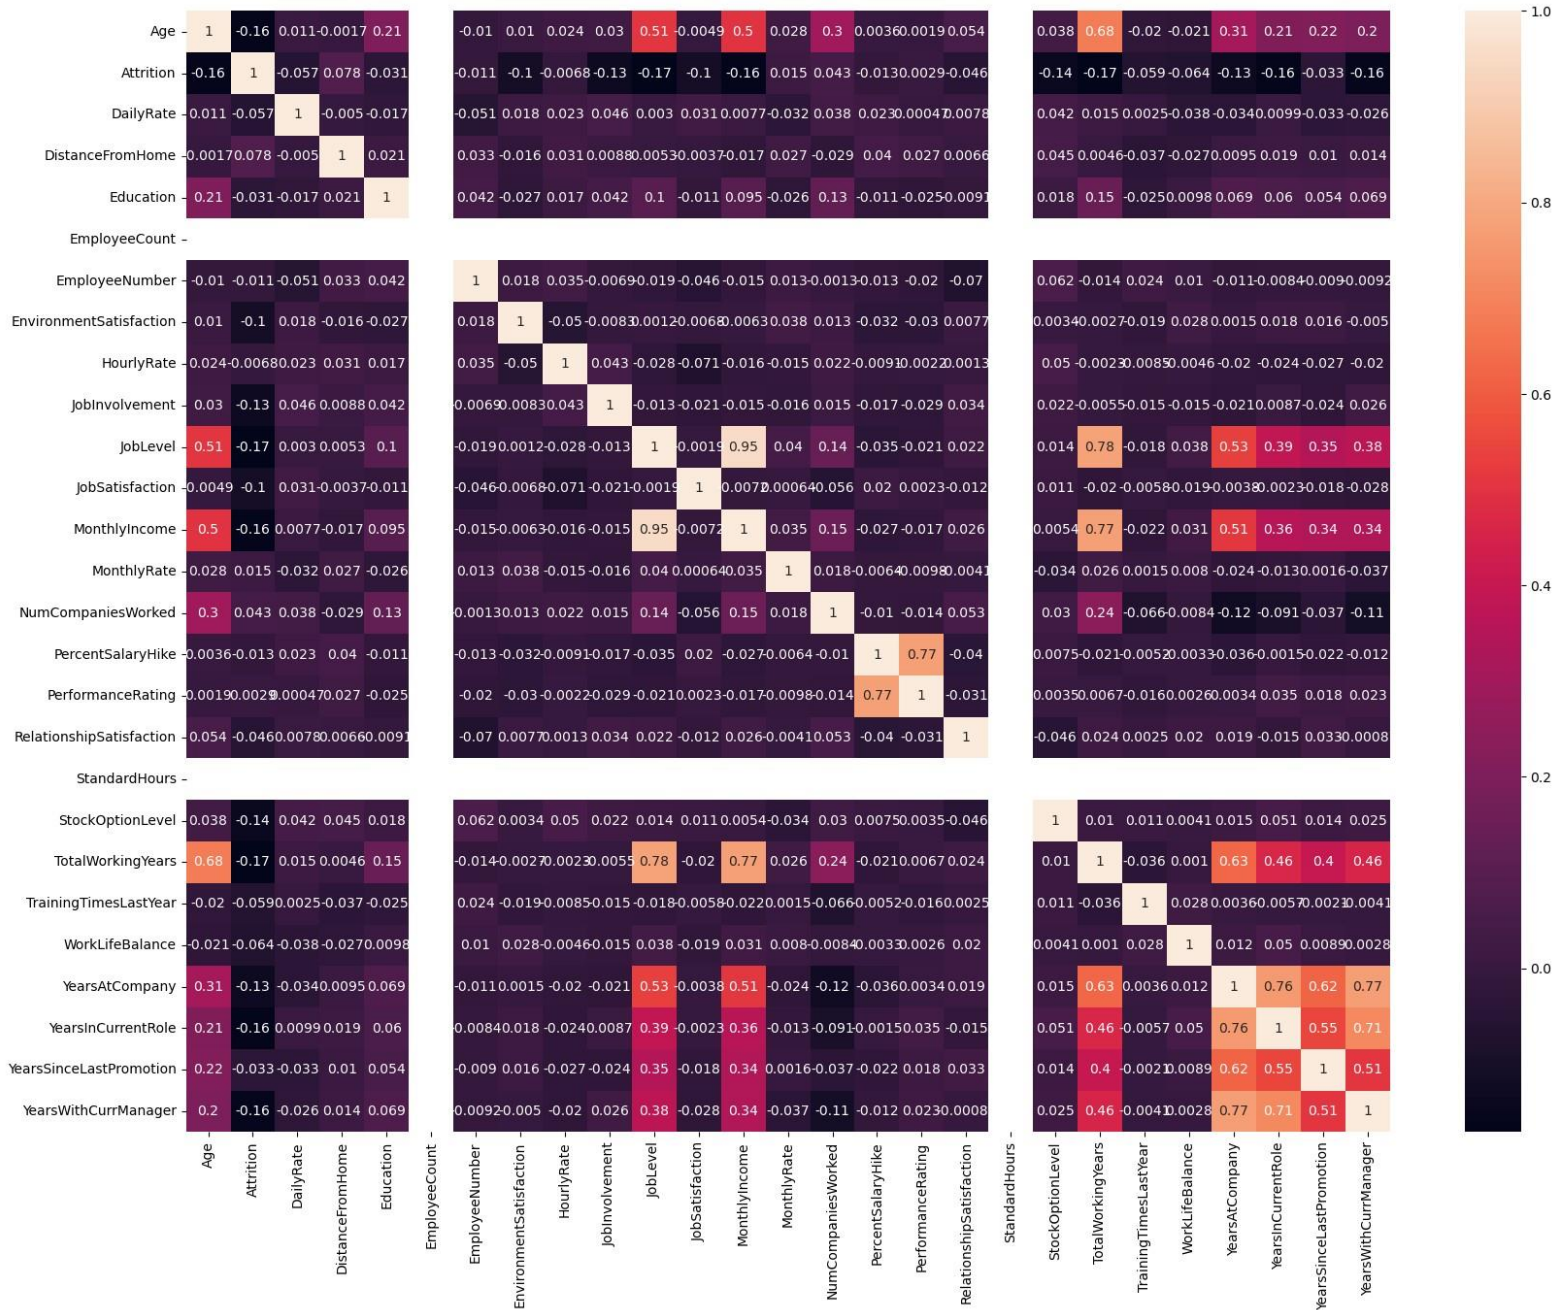
5 rows × 35 columns

Department DistanceFromHome Education EducationField EmployeeCount EmployeeNumber ... Re

4.Data Visualization.


```
In [14]: correlation_matrix = dataset.corr()
plt.figure(figsize=(20, 15))
sns.heatmap(correlation_matrix, annot=True)
plt.show()
```

C:\Users\rajes\AppData\Local\Temp\ipykernel_24744\4044223167.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning. correlation_matrix = dataset.corr()



```
In [19]: sns.lineplot(x="Age",y="Attrition",data=dataset)
```

Out[19]:<Axes: xlabel='Age', ylabel='Attrition'>



```
In [17]: correlation_matrix = dataset.corr()
```

```
# To get the correlation of "Attrition" with other columns attrition_correlation = correlation_matrix['Attrition'].drop('Attrition')


print(attrition_correlation)

Age -0.159205
DailyRate -0.056652
DistanceFromHome 0.077924
Education -0.031373
EmployeeCount NaN
EmployeeNumber -0.010577
EnvironmentSatisfaction -0.103369
HourlyRate -0.006846
JobInvolvement -0.130016
JobLevel -0.169105
JobSatisfaction -0.103481
MonthlyIncome -0.159840
MonthlyRate 0.015170
NumCompaniesWorked 0.043494
PercentSalaryHike -0.013478
PerformanceRating 0.002889
RelationshipSatisfaction -0.045872
StandardHours NaN
StockOptionLevel -0.137145
TotalWorkingYears -0.171063
TrainingTimesLastYear -0.059478
WorkLifeBalance -0.063939
YearsAtCompany -0.134392
YearsInCurrentRole -0.160545
YearsSinceLastPromotion -0.033019
YearsWithCurrManager -0.156199
Name: Attrition, dtype: float64
C:\Users\rajes\AppData\Local\Temp\ipykernel_24744\4043424376.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated.
In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
correlation_matrix = dataset.corr()
5.Outlier Detection
```

In [15]:

```
sns.boxplot(x=dataset["MonthlyIncome"])
```

Out[15]:<Axes: xlabel='MonthlyIncome'>



Inference : It shows that MonthlyIncome has outliers

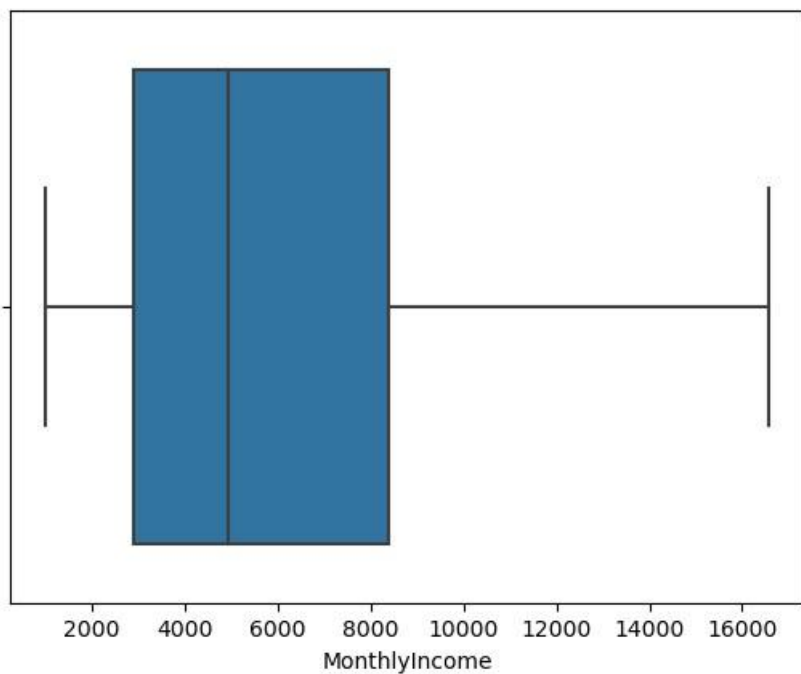
Inference : It shows that fare has outliers that are left skewed.

In [16]:

```
#So we will use flooring and capping for removing outliers
Q1 = dataset['MonthlyIncome'].quantile(0.25) Q3 = dataset['MonthlyIncome'].quantile(0.75) IQR = Q3 - Q1 whisker_width = 1.5 lower_whisker = Q1 - (whisker_width*IQR) upper_whisker = Q3 + (whisker_width*IQR)
dataset['MonthlyIncome']=np.where(dataset['MonthlyIncome']>upper_whisker,upper_whisker,np.where(dataset['MonthlyIncome']<lower_whisker,lower
```

In [17]: sns.boxplot(x=dataset["MonthlyIncome"])

Out[17]:<Axes: xlabel='MonthlyIncome'>



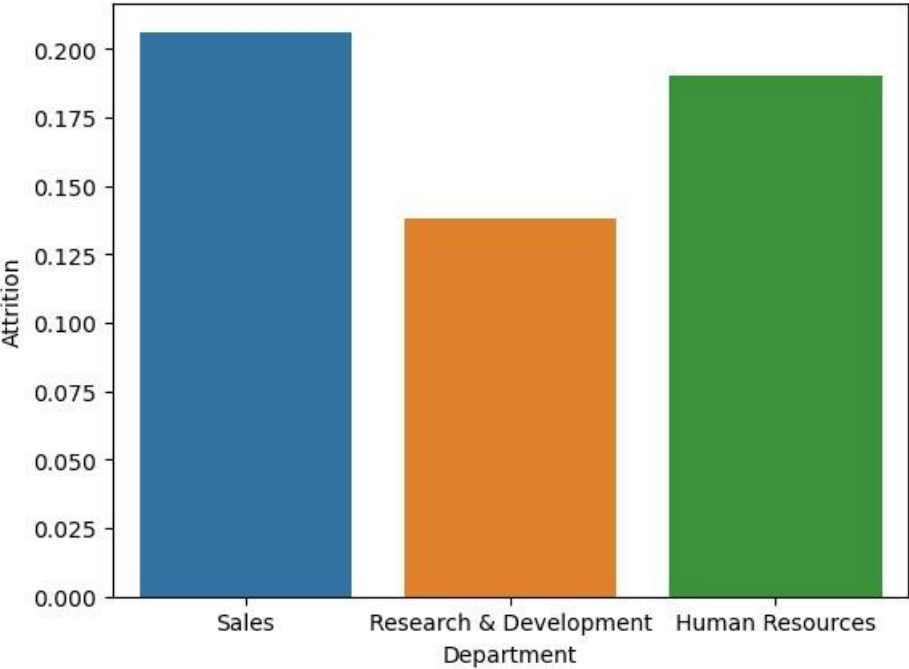
Inference: Hence We have successfully removed outliers.

```
In [21]:
sns.barplot(y=dataset['Attrition'],x=dataset['Department'],ci=0)

C:\Users\rajes\AppData\Local\Temp\ipykernel_25492\729302506.py:1: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=(‘ci’, 0)` for the same effect.

```
sns.barplot(y=dataset['Attrition'],x=dataset['Department'],ci=0) Out[21]:<Axes:
xlabel='Department', ylabel='Attrition'>
```



```
In [ ]:
Inference: It shows that Sales columns has the most attrition.
```

6.Splitting Dependent and Independent variables

```
In [22]:
#dropping unnecessary columns
dataset.drop(['Over18','EmployeeCount'],axis=1,inplace=True)

In [24]: x=dataset.drop(columns=['Attrition'])
y=dataset.iloc[:,0:1]

In [25]:
y.shape

Out[25]:(1470, 1) 7.Perform
Encoding

In [40]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
```

```
columns_to_encode = ['BusinessTravel', 'Department', 'EducationField','Gender','JobRole','MaritalStatus','OverTime']
```

```
# Create a dictionary to store label encoders and mappings
label_encoders = {}

# Perform label encoding and print mappings
for column in columns_to_encode:
    le = LabelEncoder()
    x[column] = le.fit_transform(x[column])
    label_encoders[column] = dict(zip(le.classes_, le.transform(le.classes_)))
```

```
x.head(5)
```

Out[40]:

Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeNumber	EnvironmentSatisfaction	Gender			
0	41	2		1102	2	1	2	1	1	2	0	
1	49	1		279	1	8	1	1	2	3	1	
2	37	2		1373	1	2	2	4	4	4	1	
3	33	1		1392	1	3	4	1	5	4	0	
4	27	2		591	1	2	1	3	7	1	1	

5 rows x 32 columns

In [42]:

```
# Print the mappings for each column
print("\nMappings:")
for column, mapping in label_encoders.items():
    print(f"{column} Mapping:")
    for key, value in mapping.items():
        print(f"{key}: {value}")
    print()
```

Mappings:

BusinessTravel Mapping:

Non-Travel: 0

Travel_Frequently: 1

Travel_Rarely: 2

Department Mapping:

Human Resources: 0

Research & Development: 1

Sales: 2

EducationField Mapping:

Human Resources: 0

Life Sciences: 1

Marketing: 2

Medical: 3

Other: 4

Technical Degree: 5

Gender Mapping:

Female: 0

Male: 1

JobRole Mapping:

Healthcare Representative: 0

Human Resources: 1

Laboratory Technician: 2

Manager: 3

Manufacturing Director: 4

Research Director: 5

Research Scientist: 6

Sales Executive: 7

Sales Representative: 8

MaritalStatus Mapping:

Divorced: 0

Married: 1

Single: 2

OverTime Mapping:

No: 0

Yes: 1

8.Feature Scaling

In [43]:


```
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
```

```
In [44]: x=sc.fit_transform(x)
```

```
In [45]:
x
```

```
Out[45]:array([[ 0.4463504 ,  0.59004834,  0.74252653, ..., -0.0632959 ,
                -0.67914568,  0.24583399],
               [ 1.32236521, -0.91319439, -1.2977746 , ...,  0.76499762,
                -0.36871529,  0.80654148],
               [ 0.008343 ,  0.59004834,  1.41436324, ..., -1.16768726,
                -0.67914568, -1.15593471],
               ...,
               [-1.08667552,  0.59004834, -1.60518328, ..., -0.61549158,
                -0.67914568, -0.31487349],
               [ 1.32236521, -0.91319439,  0.54667746, ...,  0.48889978,
                -0.67914568,  1.08689522],
               [-0.32016256,  0.59004834, -0.43256792, ..., -0.33939374,
                -0.36871529, -0.59522723]])
```

9.Splitting Data into Train and Test

```
In [46]:
from sklearn.model_selection import train_test_split x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
```

```
In [49]:
print(x_train.shape)
print(y_train.shape) print(x_test.shape)
print(y_test.shape)
```

```
(1029, 32)
```

```
(1029, 1)
```

```
(441, 32)
```

```
(441, 1)
```

```
In [50]:
x.shape
```

```
Out[50]:(1470, 32)
```

Model Building

1.Logistic Regression

```
In [57]:
from sklearn.linear_model import LogisticRegression from sklearn.metrics import
accuracy_score, classification_report, confusion_matrix
```

```
# Create and train the Logistic Regression model
logistic_regression_model = LogisticRegression()
logistic_regression_model.fit(x_train, y_train)
```

```
# Make predictions y_pred_lr =
logistic_regression_model.predict(x_test)
```

```
# Calculate performance metrics accuracy_lr =
accuracy_score(y_test, y_pred_lr) confusion_matrix_lr =
confusion_matrix(y_test, y_pred_lr) classification_report_lr =
classification_report(y_test, y_pred_lr)
```

```
# Print the metrics print("Logistic
Regression Metrics:")
print(f"Accuracy: {accuracy_lr}")
print("Confusion Matrix:")
print(confusion_matrix_lr)
print("Classification Report:")
print(classification_report_lr)
```



```
C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
Logistic Regression Metrics: Accuracy:
0.11337868480725624
Confusion Matrix:
[[1 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0] [1
 0 0 ... 0 0 0]
 ...
 [0 0 0 ... 1 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 1 0]]
Classification Report:
      precision    recall  f1-score   support

18      0.50      0.50      0.50         2
19      0.00      0.00      0.00         3
20      0.00      0.00      0.00         3
21      0.17      0.50      0.25         2
22      0.10      0.25      0.14         4
23      0.00      0.00      0.00         3
24      0.08      0.11      0.10         9
25      0.40      0.17      0.24        12
26      0.08      0.11      0.09         9
27      0.06      0.06      0.06        16
28      0.23      0.27      0.25        11
29      0.14      0.26      0.18        19
30      0.12      0.07      0.09        14
31      0.11      0.04      0.06        27
32      0.14      0.11      0.12        18
33      0.14      0.13      0.14        15
34      0.25      0.17      0.20        30
35      0.09      0.13      0.11        23
36      0.12      0.14      0.13        21
37      0.25      0.15      0.19        13
38      0.08      0.05      0.06        19
39      0.15      0.22      0.18         9
40      0.07      0.07      0.07        14
41      0.00      0.00      0.00        11
42      0.12      0.08      0.10        12
43      0.09      0.07      0.08        14
44      0.07      0.11      0.09         9
45      0.00      0.00      0.00        10
46      0.11      0.11      0.11         9
47      0.00      0.00      0.00         8
48      0.00      0.00      0.00         5
49      0.00      0.00      0.00         6
50      0.09      0.08      0.08        13
51      0.00      0.00      0.00         5
52      0.00      0.00      0.00         6
53      0.14      0.14      0.14         7
54      0.25      0.50      0.33         4
55      0.33      0.20      0.25        10
56      0.00      0.00      0.00         5
57      0.00      0.00      0.00         1
58      0.17      0.20      0.18         5
59      0.00      0.00      0.00         1      60      0.00      0.00      0.00         4

      accuracy          0.11      441  macro avg
0.11      0.12      0.11      441 weighted avg      0.12
0.11      0.11      441
```

C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in: <https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression n_iter_i =
_check_optimize_result(

C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

2.Decision Tree

In [53]:

```
from sklearn.tree import DecisionTreeClassifier

# Create and train the Decision Tree model decision_tree_model =
DecisionTreeClassifier() decision_tree_model.fit(x_train, y_train)

# Make predictions y_pred_dt =
decision_tree_model.predict(x_test)

# Calculate performance metrics accuracy_dt =
accuracy_score(y_test, y_pred_dt) confusion_matrix_dt =
confusion_matrix(y_test, y_pred_dt) classification_report_dt =
classification_report(y_test, y_pred_dt)

# Print the metrics print("Decision
Tree Metrics:") print(f"Accuracy:
{accuracy_dt}") print("Confusion
Matrix:")
print(confusion_matrix_dt)
print("Classification Report:")
print(classification_report_dt)
```

Decision Tree Metrics:
Accuracy: 0.9909297052154195
Confusion Matrix:
[[2 0 0 ... 0 0 0]
[0 3 0 ... 0 0 0] [0
0 3 ... 0 0 0]
...
[0 0 0 ... 5 0 0]
[0 0 0 ... 0 1 0]
[0 0 0 ... 0 4 0]]
Classification Report:

	precision	recall	f1-score	support					
18	1.00	1.00	1.00	2					
19	1.00	1.00	1.00	3					
20	1.00	1.00	1.00	3					
21	1.00	1.00	1.00	2					
22	1.00	1.00	1.00	4					
23	1.00	1.00	1.00	3					
24	1.00	1.00	1.00	9					
25	1.00	1.00	1.00	12					
26	1.00	1.00	1.00	9					
27	1.00	1.00	1.00	16					
28	1.00	1.00	1.00	11					
29	1.00	1.00	1.00	19					
30	1.00	1.00	1.00	14					
31	1.00	1.00	1.00	27					
32	1.00	1.00	1.00	18					
33	1.00	1.00	1.00	15					
34	1.00	1.00	1.00	30					
35	1.00	1.00	1.00	23					
36	1.00	1.00	1.00	21					
37	1.00	1.00	1.00	13					
38	1.00	1.00	1.00	19					
39	1.00	1.00	1.00	9					
40	1.00	1.00	1.00	14					
41	1.00	1.00	1.00	11					
42	1.00	1.00	1.00	12					
43	1.00	1.00	1.00	14					
44	1.00	1.00	1.00	9					
45	1.00	1.00	1.00	10					
46	1.00	1.00	1.00	9					
47	1.00	1.00	1.00	8					
48	1.00	1.00	1.00	5					
49	1.00	1.00	1.00	6					
50	1.00	1.00	1.00	13					
51	1.00	1.00	1.00	5					
52	1.00	1.00	1.00	6					
53	1.00	1.00	1.00	7					
54	1.00	1.00	1.00	4					
55	1.00	1.00	1.00	10					
56	1.00	1.00	1.00	5					
57	1.00	1.00	1.00	1					
58	1.00	1.00	1.00	5					
59	0.20	1.00	0.33	1	60	0.00	0.00	0.00	4
accuracy			0.99	441	macro avg				
0.96	0.98	0.96	441	weighted avg	0.99				
0.99	0.99	441							

C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior. _warn_prf(average, modifier, msg_start, len(result))
C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior. _warn_prf(average, modifier, msg_start, len(result))
C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior. _warn_prf(average, modifier, msg_start, len(result))

3.Random Forest

In [58]: **from** sklearn.ensemble **import** RandomForestClassifier

```
# Create and train the Random Forest model random_forest_model  
= RandomForestClassifier() random_forest_model.fit(x_train,  
y_train)
```

```
# Make predictions y_pred_rf =  
random_forest_model.predict(x_test) # Calculate  
performance metrics accuracy_rf =  
accuracy_score(y_test, y_pred_rf)  
confusion_matrix_rf = confusion_matrix(y_test,  
y_pred_rf) classification_report_rf =  
classification_report(y_test, y_pred_rf)
```

```
# Print the metrics print("Random  
Forest Metrics:") print(f"Accuracy:  
{accuracy_rf}") print("Confusion  
Matrix:")  
print(confusion_matrix_rf)  
print("Classification Report:")  
print(classification_report_rf)
```

C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\base.py:1151: DataConversionWarning: A column-vector y was passed when a 1d array was expected.

Please change the shape of y to (n_samples,), for example using ravel().

return fit_method(estimator, *args, **kwargs)

Random Forest Metrics: Accuracy:

0.6077097505668935

Confusion Matrix:

[[2 0 0 ... 0 0 0]
[0 1 2 ... 0 0 0] [0
0 3 ... 0 0 0]
...
[0 0 0 ... 0 0 0]
[0 0 0 ... 0 0 0]
[0 0 0 ... 0 0 0]]

Classification Report:

	precision	recall	f1-score	support					
18	1.00	1.00	1.00	2					
19	1.00	0.33	0.50	3					
20	0.60	1.00	0.75	3					
21	0.50	1.00	0.67	2					
22	0.33	0.25	0.29	4					
23	0.00	0.00	0.00	3					
24	0.33	0.33	0.33	9					
25	0.62	0.42	0.50	12					
26	0.56	0.56	0.56	9					
27	0.69	0.56	0.62	16					
28	0.80	0.73	0.76	11					
29	0.68	0.79	0.73	19					
30	0.65	0.93	0.76	14					
31	0.90	1.00	0.95	27					
32	1.00	0.89	0.94	18					
33	1.00	0.93	0.97	15					
34	0.94	1.00	0.97	30					
35	0.82	1.00	0.90	23					
36	0.84	1.00	0.91	21					
37	0.65	0.85	0.73	13					
38	0.64	0.84	0.73	19					
39	0.44	0.44	0.44	9					
40	0.43	0.71	0.54	14					
41	0.33	0.27	0.30	11					
42	0.29	0.50	0.36	12					
43	0.20	0.07	0.11	14					
44	0.00	0.00	0.00	9					
45	0.27	0.40	0.32	10					
46	0.27	0.44	0.33	9					
47	0.00	0.00	0.00	8					
48	0.00	0.00	0.00	5					
49	0.20	0.17	0.18	6					
50	0.50	0.46	0.48	13					
51	0.20	0.20	0.20	5					
52	0.00	0.00	0.00	6					
53	1.00	0.29	0.44	7					
54	0.00	0.00	0.00	4					
55	0.25	0.10	0.14	10					
56	0.00	0.00	0.00	5					
57	0.00	0.00	0.00	1					
58	0.00	0.00	0.00	5					
59	0.00	0.00	0.00	1	60	0.00	0.00	0.00	4
accuracy			0.61	441	macro avg				
0.44	0.45	0.43	441	weighted avg	0.57				
0.61	0.58	441							

C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\rajes\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))